

```
In [1]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier
from sklearn.tree import plot_tree
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.metrics import silhouette_score
import seaborn as sns
```

```
In [2]: pd.set_option('display.max_columns', None)
```

HW:

The data set includes the churn of customers of a telecommunications company. The task is to create segments from customers based on their characteristics using the KMeans algorithm.

Do not use the following variables for grouping:

- churn?: has the customer dropped out?
- Contract_date: contract conclusion time
- Cust_ID: customer ID

```
In [3]: file_path = "/content/telco_sampled.csv"
df = pd.read_csv(file_path, sep = ';')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Contract_date	Package	Gender	Age	Marital_Status	Living_Condition	Graduation	Job_Type	Income	Peak_minute_09	Week
0	9/20/04 12:00 AM	PACK_B	Male	42.0	Married	Owner	University	Leader	15_30k	0.55	
1	2/12/05 12:00 AM	PACK_B	Female	53.0	Married	Owner	University	Public_Employee	Below_15k	11.32	
2	10/19/04 12:00 AM	PACK_X	Male	43.0	Married	Owner	Highschool	Executive	30_60k	78.05	
3	10/31/04 12:00 AM	PACK_B	Male	32.0	Married	Owner	Highschool	Labourer	15_30k	0.08	
4	11/19/04 12:00 AM	PACK_B	Female	31.0	Married	Owner	Highschool	Public_Employee	30_60k	20.68	

```
In [5]: df['churn?'].value_counts()
```

```
Out[5]:
```

	count
churn?	
0	1224
1	341

dtype: int64

1. Subtask: (data preparation)

Use all variables except for the three variables above when creating the clusters. Perform data preparation so that the variables are input to the model in the appropriate form.

(hint: categorical variables, missing values, scaling, etc.)

```
In [6]: # Dropping the columns not required for clustering
df = df.drop(columns=['churn?', 'Contract_date', 'Cust_ID'])
```

```
In [7]: # Temporarily convert 'Age' to object type to avoid scaling
df['Age'] = df['Age'].astype('object')
```

```
In [8]: # Handling categorical variables with OneHotEncoding and missing values
categorical_features = df.select_dtypes(include=['object']).columns.tolist()
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
```

```
In [9]: # Now 'Age' will be in categorical features
print("Categorical Features:", categorical_features)
print("Numerical Features:", numerical_features)

Categorical Features: ['Package', 'Gender', 'Age', 'Marital_Status', 'Living_Condition', 'Graduation', 'Job_Type',
'Income']
Numerical Features: ['Peak_minute_09', 'Weekend_minute_09', 'Offpeak_minute_09', 'Offpeak_nr_09', 'Peak_nr_09', 'Weekend_nr_09', 'Selfnet_minute_09', 'Fixed_minute_09', 'Othermob_minute_09', 'Voicemail_nr_09', 'Voicemail_minute_09', 'SMS_09', 'Peak_minute_10', 'Weekend_minute_10', 'Offpeak_minute_10', 'Offpeak_nr_10', 'Peak_nr_10', 'Weekend_nr_10', 'Selfnet_minute_10', 'Fixed_minute_10', 'Othermob_minute_10', 'Voicemail_nr_10', 'Voicemail_minute_10', 'SMS_10', 'Peak_minute_11', 'Weekend_minute_11', 'Offpeak_minute_11', 'Offpeak_nr_11', 'Peak_nr_11', 'Weekend_nr_11', 'Selfnet_minute_11', 'Fixed_minute_11', 'Othermob_minute_11', 'Voicemail_nr_11', 'Voicemail_minute_11', 'SMS_11', 'Peak_minute_12', 'Weekend_minute_12', 'Offpeak_minute_12', 'Offpeak_nr_12', 'Peak_nr_12', 'Weekend_nr_12', 'Selfnet_minute_12', 'Fixed_minute_12', 'Othermob_minute_12', 'Voicemail_nr_12', 'Voicemail_minute_12', 'SMS_12']
```

```
In [10]: # Apply scaling only to numerical columns
scaler = StandardScaler()
df[numerical_features] = scaler.fit_transform(df[numerical_features])
```

```
In [12]: df.head()
```

```
Out[12]:
```

	Package	Gender	Age	Marital_Status	Living_Condition	Graduation	Job_Type	Income	Peak_minute_09	Weekend_minute_09
0	PACK_B	Male	42.0	Married	Owner	University	Leader	15_30k	-0.490696	-0.521675
1	PACK_B	Female	53.0	Married	Owner	University	Public_Employee	Below_15k	-0.290959	-0.302073
2	PACK_X	Male	43.0	Married	Owner	Highschool	Executive	30_60k	0.946595	-0.394482
3	PACK_B	Male	32.0	Married	Owner	Highschool	Labourer	15_30k	-0.499413	-0.531513
4	PACK_B	Female	31.0	Married	Owner	Highschool	Public_Employee	30_60k	-0.117372	-0.044174

```
In [11]: # Pipelines for numerical and categorical data processing
numerical_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='mean')),
    ('scaler', StandardScaler())
])

categorical_pipeline = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('encoder', OneHotEncoder(drop='first'))
])

# Combine pipelines using ColumnTransformer
preprocessor = ColumnTransformer(transformers=[
    ('num', numerical_pipeline, numerical_features),
    ('cat', categorical_pipeline, categorical_features)
])
```

```
In [13]: # Preprocess the data
X_preprocessed = preprocessor.fit_transform(df)
```

```
In [14]: df['Age'] = df['Age'].astype('int64')
```

```
In [16]: X_preprocessed_with_age = np.hstack((X_preprocessed, df[['Age']].values))
```

```
In [17]: # Display shape and preview of preprocessed data
print("Preprocessed Data Shape:", X_preprocessed.shape)
print(pd.DataFrame(X_preprocessed).head())
```

	0	1	2	3	4	5	6	\		
0	-0.490696	-0.521675	-0.459272	-0.590807	-0.567464	-0.593901	-0.439902			
1	-0.290959	-0.302073	-0.312966	0.204770	0.057757	0.073678	-0.337945			
2	0.946595	-0.394482	-0.282573	-0.437812	1.236744	-0.297199	-0.112668			
3	-0.499413	-0.531513	-0.459272	-0.590807	-0.585328	-0.630988	-0.452929			
4	-0.117372	-0.044174	0.238093	0.908549	-0.067288	0.333292	0.203970			
	7	8	9	10	11	12	13	\		
0	-0.392252	-0.455270	-0.587854	-0.566961	-0.400519	-0.499263	-0.523609			
1	0.339403	-0.318874	-0.375662	-0.379413	0.822817	0.041093	-0.363092			
2	-0.134075	0.459603	0.744239	0.230140	-0.379427	0.011658	-0.051483			
3	-0.392252	-0.455270	-0.611430	-0.573637	-0.400519	-0.499263	-0.523609			
4	-0.343901	-0.145011	0.308067	0.029330	0.147873	0.385491	0.220606			
	14	15	16	17	18	19	20	\		
0	-0.452379	-0.565336	-0.591148	-0.608457	-0.487639	-0.277618	-0.448946			
1	-0.071564	1.266374	0.970777	0.455122	-0.266674	0.493279	-0.131202			
2	-0.387259	-0.248309	0.220242	-0.076668	-0.145313	-0.176438	-0.119198			
3	-0.452379	-0.565336	-0.591148	-0.608457	-0.487639	-0.277618	-0.448946			
4	0.317670	0.808446	0.544798	0.322174	0.821386	-0.185472	-0.080363			
	21	22	23	24	25	26	27	\		
0	-0.601178	-0.568592	-0.392033	-0.473401	-0.477615	-0.442899	-0.550587			
1	-0.222053	-0.223253	2.288882	0.023522	-0.477615	-0.247176	1.402267			
2	0.222059	-0.175703	-0.392033	0.610884	-0.436604	-0.397831	-0.282548			
3	-0.601178	-0.568592	-0.392033	-0.473401	-0.477615	-0.442899	-0.550587			
4	0.662572	0.334952	0.089694	0.144957	0.435284	0.287716	0.713025			
	28	29	30	31	32	33	34	\		
0	-0.534268	-0.551492	-0.452852	-0.309751	-0.433061	-0.575227	-0.540116			
1	0.371554	0.154393	-0.417322	0.743390	-0.180469	-0.089204	-0.222023			
2	0.648847	-0.410315	-0.324128	0.810843	0.118437	0.423819	0.022850			
3	-0.534268	-0.551492	-0.452852	-0.309751	-0.433061	-0.575227	-0.540116			
4	0.353068	1.072043	0.533448	-0.264782	0.041806	0.815338	0.297378			
	35	36	37	38	39	40	41	\		
0	-0.353420	-0.472188	-0.361515	-0.449857	-0.577982	-0.609884	-0.596913			
1	1.391981	-0.250612	-0.272427	-0.339479	0.304567	0.210353	0.322726			
2	-0.332392	-0.287325	-0.147455	-0.340768	-0.472076	-0.158754	0.065227			
3	-0.353420	-0.472188	-0.361515	-0.449857	-0.577982	-0.609884	-0.596913			
4	0.088187	0.448977	0.526895	0.309636	1.857853	0.989578	1.132007			
	42	43	44	45	46	47	48	49	50	\
0	-0.429383	-0.245480	-0.454958	-0.588222	-0.543702	-0.435443	1.0	0.0	0.0	
1	-0.412159	0.112460	-0.266448	-0.311383	-0.357744	0.975686	1.0	0.0	0.0	
2	-0.374192	-0.223542	-0.315431	-0.179556	-0.336739	-0.178874	0.0	0.0	0.0	
3	-0.429383	-0.245480	-0.454958	-0.588222	-0.543702	-0.435443	1.0	0.0	0.0	
4	0.888704	-0.122510	0.0059							

	126	127	128	129
0	0.0	0.0	0.0	0.0
1	0.0	0.0	1.0	0.0
2	0.0	1.0	0.0	0.0
3	0.0	0.0	0.0	0.0
4	0.0	1.0	0.0	0.0

2. Subtask: (clustering)


Find the optimal k value for the KMeans algorithm using the variables prepared in the previous task. Then group the customers.

```
In [18]: # Function to apply only the Elbow Method
def elbow_method(X):
    wcss = [] # Within-cluster sum of squares (WCSS) for the elbow method
    K_values = range(1, 11) # Trying out K values from 1 to 10

    # Loop to calculate WCSS for each K value
    for k in K_values:
        kmeans = KMeans(n_clusters=k, random_state=42)
        kmeans.fit(X)
        wcss.append(kmeans.inertia_)

    # Plotting the elbow graph
    plt.figure(figsize=(8, 5))
    plt.plot(K_values, wcss, 'bo-')
    plt.title('Elbow Method for Optimal K')
    plt.xlabel('Number of clusters (K)')
    plt.ylabel('WCSS (Inertia)')
    plt.show()

# Assuming X_preprocessed is already available from previous steps
elbow_method(X_preprocessed)
```



```
In [19]: # The optimal k is 4
# Apply KMeans clustering with the optimal number of clusters
k = 4
kmeans = KMeans(n_clusters=k, random_state=42)
df['Cluster'] = kmeans.fit_predict(X_preprocessed)
```

Let's generate graphs for different customer segments based on their characteristics.

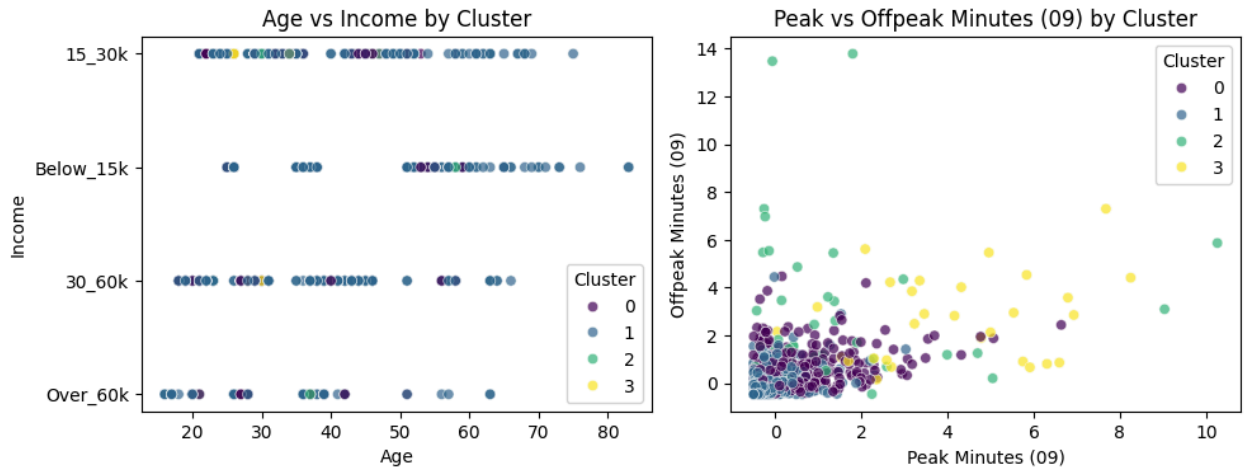
```
In [20]: # Create scatter plots for key feature relationships
plt.figure(figsize=(10, 7))

# Scatter plot of Age vs Income colored by Cluster
plt.subplot(2, 2, 1)
sns.scatterplot(data=df, x='Age', y='Income', hue='Cluster', palette='viridis', alpha=0.7)
plt.title('Age vs Income by Cluster')
plt.xlabel('Age')
plt.ylabel('Income')

# Scatter plot of Peak_minute_09 vs Offpeak_minute_09 colored by Cluster
```

```
plt.subplot(2, 2, 2)
sns.scatterplot(data=df, x='Peak_minute_09', y='Offpeak_minute_09', hue='Cluster', palette='viridis', alpha=0.7)
plt.title('Peak vs Offpeak Minutes (09) by Cluster')
plt.xlabel('Peak Minutes (09)')
plt.ylabel('Offpeak Minutes (09)')

plt.tight_layout()
plt.show()
```



```
In [22]: # Count plot for Gender distribution across clusters
plt.figure(figsize=(8, 4))
sns.countplot(data=df, x='Gender', hue='Cluster', palette='viridis')
plt.title('Gender Distribution by Cluster')
plt.xlabel('Gender')
plt.ylabel('Count')
plt.legend(title='Cluster')
plt.grid(axis='y')
plt.show()
```

/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to pass a length-1 tuple to get_group in a future version of pandas. Pass `(name,)` instead of `name` to silence this warning.

data_subset = grouped_data.get_group(pd_key)

/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to pass a length-1 tuple to get_group in a future version of pandas. Pass `(name,)` instead of `name` to silence this warning.


data_subset = grouped_data.get_group(pd_key)

/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to pass a length-1 tuple to get_group in a future version of pandas. Pass `(name,)` instead of `name` to silence this warning.

data_subset = grouped_data.get_group(pd_key)

/usr/local/lib/python3.10/dist-packages/seaborn/_base.py:949: FutureWarning: When grouping with a length-1 list-like, you will need to pass a length-1 tuple to get_group in a future version of pandas. Pass `(name,)` instead of `name` to silence this warning.

data_subset = grouped_data.get_group(pd_key)



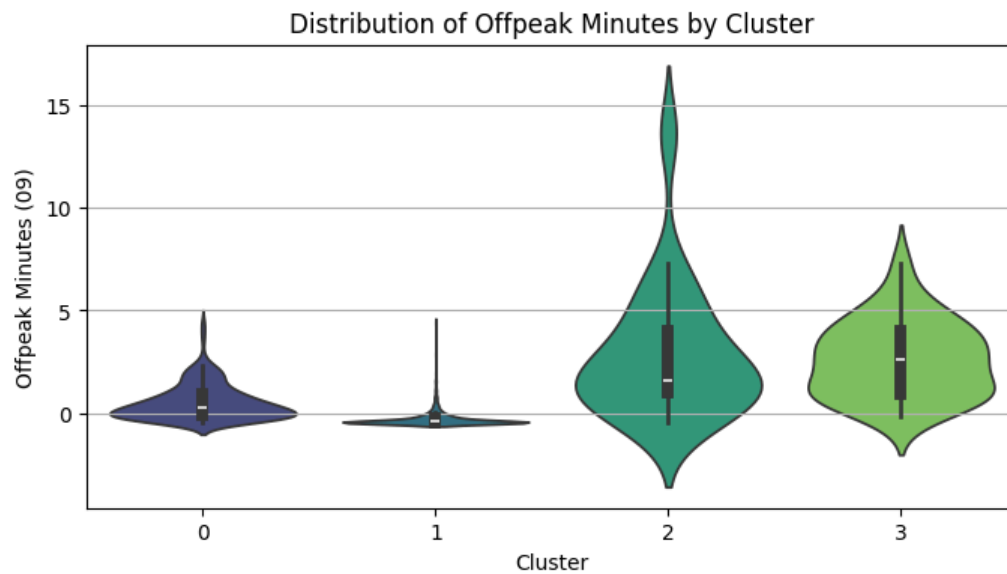
```
In [23]: # Violin plot for Offpeak Minutes by Cluster
plt.figure(figsize=(8, 4))
sns.violinplot(data=df, x='Cluster', y='Offpeak_minute_09', palette='viridis')
plt.title('Distribution of Offpeak Minutes by Cluster')
```

```
plt.xlabel('Cluster')
plt.ylabel('Offpeak Minutes (09)')
plt.grid(axis='y')
plt.show()
```

<ipython-input-23-368e32be3b7c>:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.violinplot(data=df, x='Cluster', y='Offpeak_minute_09', palette='viridis')
```



3. Subtask: (explanation of clusters / conclusions)

Try to find an explanation of what characterizes each group and what characteristics caused each customer to be in the given cluster.

```
In [24]: # Analyze cluster characteristics
numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns.tolist()
cluster_summary = df.groupby('Cluster')[numerical_columns].mean()
# Display cluster summary
print(cluster_summary)
```

	Age	Peak_minute_09	Weekend_minute_09	Offpeak_minute_09	\
Cluster					
0	31.557554	0.623515	0.648169	0.542985	
1	34.912151	-0.288483	-0.291217	-0.282465	
2	30.564103	1.466402	2.179957	2.933920	
3	31.000000	4.028172	2.983102	2.622342	

	Offpeak_nr_09	Peak_nr_09	Weekend_nr_09	Selfnet_minute_09	\
Cluster					
0	0.679826	0.728985	0.697765	0.583021	
1	-0.304110	-0.302754	-0.295190	-0.271121	
2	1.330660	0.761302	0.971390	2.240069	
3	4.317289	4.546841	4.255933	2.692765	

	Fixed_minute_09	Othermob_minute_09	Voicemail_nr_09	\
Cluster				
0	0.469722	0.529366	0.713479	
1	-0.175707	-0.276530	-0.306159	
2	1.428487	2.072138	1.078806	
3	0.923905	3.627901	4.568326	

	Voicemail_minute_09	SMS_09	Peak_minute_10	Weekend_minute_10	\
Cluster					
0	0.709356	0.513341	0.631212	0.686603	
1	-0.337058	-0.198136	-0.299180	-0.324594	
2	2.491617	0.534560	1.717162	2.685880	
3	3.955140	2.592436	4.065163	3.324369	

	Offpeak_minute_10	Offpeak_nr_10	Peak_nr_10	Weekend_nr_10	\
Cluster					
0	0.601369	0.754595	0.776540	0.788327	
1	-0.296855	-0.323040	-0.318153	-0.330692	
2	3.020337	1.501208	0.895880	1.416433	
3	2.553179	4.171277	4.556444	4.279573	

	Selfnet_minute_10	Fixed_minute_10	Othermob_minute_10	\
Cluster				
0	0.649073	0.346518	0.567723	
1	-0.302837	-0.154118	-0.292401	
2	2.708910	1.857020	2.152563	
3	2.758863	0.631981	3.812239	

	Voicemail_nr_10	Voicemail_minute_10	SMS_10	Peak_minute_11	\
Cluster					
0	0.773335	0.732817	0.585660	0.577917	
1	-0.332966	-0.359185	-0.207297	-0.291288	
2	1.405467	2.959405	0.525772	1.718338	
3	4.681107	3.944900	2.305638	4.237088	

	Weekend_minute_11	Offpeak_minute_11	Offpeak_nr_11	Peak_nr_11	\
Cluster					
0	0.569108	0.484439	0.737818	0.710622	
1	-0.305210	-0.293397	-0.327976	-0.305103	
2	2.390085	3.406594	1.565251	0.866889	
3	4.010681	2.994212	4.443870	4.675133	

	Weekend_nr_11	Selfnet_minute_11	Fixed_minute_11	\
Cluster				
0	0.734462	0.479425	0.320243	
1	-0.320051	-0.281590	-0.168695	
2	1.188785	2.990417	2.268665	
3	4.642643	3.102333	0.932162	

	Othermob_minute_11	Voicemail_nr_11	Voicemail_minute_11	SMS_11	\
Cluster					
0	0.555094	0.729924	0.649074	0.611566	
1	-0.292224	-0.329638	-0.344725	-0.210085	
2	1.747592	1.410299	2.865366	0.363181	
3	4.448558	4.950998	4.402848	2.390154	

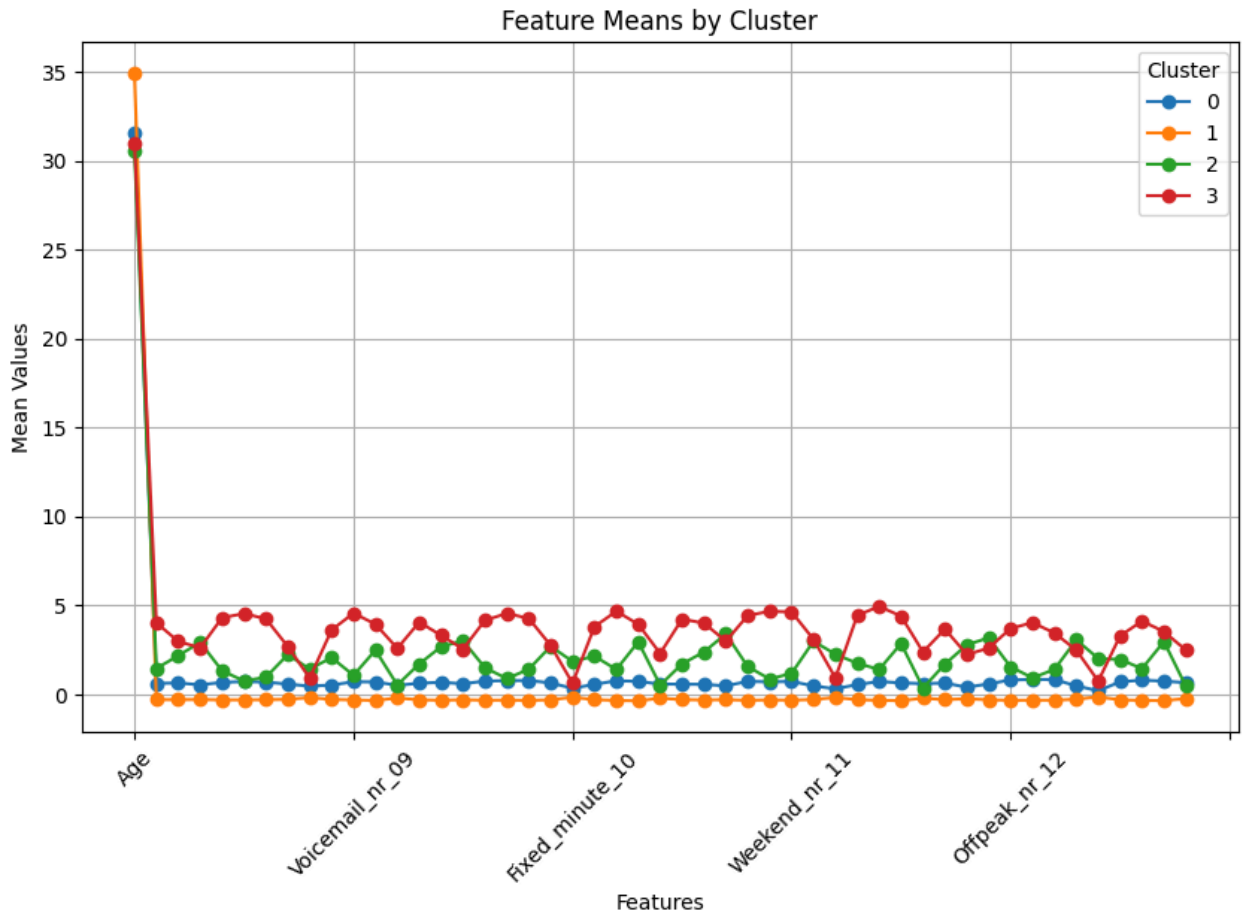
	Peak_minute_12	Weekend_minute_12	Offpeak_minute_12	Offpeak_nr_12	\
Cluster					
0	0.640959	0.407556	0.595196	0.817334	
1	-0.291453	-0.237706	-0.299381	-0.327767	
2	1.704663	2.791620	3.186739	1.550092	
3	3.677387	2.245081	2.602870	3.718265	

	Peak_nr_12	Weekend_nr_12	Selfnet_minute_12	Fixed_minute_12	\
Cluster					
0	0.845152	0.839311	0.483180	0.208477	
1	-0.322296	-0.321867	-0.267422	-0.130949	
2	0.935947	1.408371	3.068274	2.024555	
3	4.036760	3.459306	2.493391	0.752731	

Othermob_minute_12 Voicemail_nr_12 Voicemail_minute_12 SMS_12


Cluster				
0	0.712576	0.803938	0.750995	0.664358
1	-0.304043	-0.331731	-0.345962	-0.228043
2	1.965588	1.454773	2.977059	0.465837
3	3.295514	4.127263	3.532892	2.496569

```
In [25]: # Visualizing the mean values for each feature across clusters
cluster_summary.T.plot(marker='o', figsize=(10, 6))
plt.xticks(rotation=45)
plt.title('Feature Means by Cluster')
plt.xlabel('Features')
plt.ylabel('Mean Values')
plt.legend(title='Cluster')
plt.grid()
plt.show()
```



```
In [36]: # Exclude the Age feature for better visualization
cluster_summary_excluding_age = cluster_summary.drop(columns=['Age'])

# Visualizing the mean values for each feature across clusters (excluding Age)
cluster_summary_excluding_age.T.plot(marker='o', figsize=(13, 6))
plt.xticks(rotation=45)
plt.title('Feature Means by Cluster (Excluding Age)')
plt.xlabel('Features')
plt.ylabel('Mean Values')
plt.legend(title='Cluster')
plt.grid()
plt.show()
```

```
In [33]: # Generate characteristics for each cluster
for cluster in range(k):
    print(f"\nCluster {cluster} Characteristics:")
    characteristics = cluster_summary.loc[cluster]

    # Print Age and other characteristics
    print(f"- Average Age: {characteristics['Age']:.1f} years")
    print(f"- Peak Minutes: {characteristics['Peak_minute_09']:.2f}")
    print(f"- Off-Peak Minutes: {characteristics['Offpeak_minute_09']:.2f}")
    print(f"- Weekend Minutes: {characteristics['Weekend_minute_09']:.2f}")
    print(f"- Selfnet Minutes: {characteristics['Selfnet_minute_09']:.2f}")
    print(f"- Othermob Minutes: {characteristics['Othermob_minute_09']:.2f}")
```

Cluster 0 Characteristics:

- Average Age: 31.6 years
- Peak Minutes: 0.62
- Off-Peak Minutes: 0.54
- Weekend Minutes: 0.65
- Selfnet Minutes: 0.58
- Othermob Minutes: 0.53

Cluster 1 Characteristics:

- Average Age: 34.9 years
- Peak Minutes: -0.29
- Off-Peak Minutes: -0.28
- Weekend Minutes: -0.29
- Selfnet Minutes: -0.27
- Othermob Minutes: -0.28

Cluster 2 Characteristics:

- Average Age: 30.6 years
- Peak Minutes: 1.47
- Off-Peak Minutes: 2.93
- Weekend Minutes: 2.18
- Selfnet Minutes: 2.24
- Othermob Minutes: 2.07

Cluster 3 Characteristics:

- Average Age: 31.0 years
- Peak Minutes: 4.03
- Off-Peak Minutes: 2.62
- Weekend Minutes: 2.98
- Selfnet Minutes: 2.69
- Othermob Minutes: 3.63

- Cluster 0: Moderately active users with balanced usage of peak and off-peak minutes. Average age is 31.6 years, likely representing young professionals with steady communication needs.
- Cluster 1: Low usage of services, indicated by negative values for peak and off-peak minutes. Average age is 34.9 years, suggesting infrequent users who may prefer alternative communication methods. Marketing efforts should aim to boost engagement.

- Cluster 2: Heavy users of mobile services, with high usage across all metrics. Average age is 30.6 years, indicating active individuals, possibly business users. Loyalty programs could help retain this segment.
- Cluster 3: High usage, particularly during peak hours, with an average age of 31.0 years. This group may consist of professionals who rely on their devices during work hours. Targeted promotions for high usage plans could be beneficial.